

ISW2 Project A.A. 2019-2020

Andrea Graziani (0273395)
andrea.graziani93@outlook.it
Università degli Studi di Roma "Tor Vergata"
Rome, Italy

KEYWORDS

Testing, Equivalence Class Partitioning, Boundary Value Analysis, Mutation Testing

1 INTRODUCTION

In this report we will describe all results and issues about the utilization of several *testing techniques* involving following open source projects:

Apache BookKeeper™ Which is a scalable, fault tolerant and low latency storage service¹.

Apache OpenJPA™ An implementation of the Java Persistence API specification².

To be more precise, the aim of this report is to analyse 2 classes of each aforementioned project, attempting both to find if there are any errors in it and to increase our confidence in the correct functioning of the **software under test (SUT)**, although the completeness of our testing does not necessarily demonstrate that these classes are error free.

1.1 Class selection

To select classes for our testing activities, we have mainly relied on results elaborated by our **defect prediction model**, developed using **Weka**³ machine learning software, using a *random forests* classifier without *sampling* or *feature selection*.

Particularly, we have focused on classes marked as *defective* (or *buggy*) by our predictor because it is likely that they exhibit a defect observable by our test sets. As known, this approach is able to *reduce the cost of testing*, by letting to focus on specific classes, ignoring stable ones [5].

In particular, for our testing activities involving Apache BookKeeper™, we have choose following classes:

DefaultEnsemblePlacementPolicy⁴ Because it has been marked as defective by our prediction model and, as we will see later, our testing activity has confirmed that prediction.

DiskChecker⁵ Although has not been marked as defective, we have choose this class due to its relative high LOC value (294), high number of revisions (11) and high number of authors (9). Despite its presumed stability, our testing activity have found several defects inside that class.

1.2 Testing technique

In order to build our test set, we have adopted a testing technique called *equivalence class partitioning*, according to which the domain

of possible input data for each input data element is divided into **equivalence classes**; an equivalence class is a set of data values that the tester assumes are processed in the same way by the test object [6].

Equivalence class definition was been mainly based on *specifications* of our test objects using, therefore, a *black-box* approach; however, sometimes, due to lacking of detailed specifications, we have performed our analysis looking the code too, using conversely a **while-box** approach.

For any equivalence classes is important to find a **representative**. As known, testing one representative of the equivalence class is considered sufficient because it is assumed that for any other input value of the same equivalence class, the test object will show the same reaction or behaviour [6]. Besides equivalence classes for *correct* input, those for *incorrect* input values must be tested as well.

To guarantee that all test object reactions are triggered, we have combined representative values using following rules:

- (1) The representative value of all valid equivalence classes have been combined to test cases (*valid/positive test case*), meaning that all possible combinations of valid equivalence classes will be covered.
- (2) The representative value of an invalid equivalence class have been combined only with representatives of other valid equivalence classes (*invalid/negative test case*).

Moreover, we have adopted following guidelines[6]:

- (1) Test cases including boundary values combinations are preferred.
- (2) Every representative of an equivalence class appears in at least one test case.
- (3) Representatives of invalid equivalence classes should not be combined with representatives of other invalid equivalence classes.

2 APACHE BOOKKEEPER™

2.1 DefaultEnsemblePlacementPolicy class testing analysis

What are responsibilities of DefaultEnsemblePlacementPolicy class? What meant by ensemble? And what for placement policy?

According to bookkeeper specifications[1], an **ensemble** represents a group of **bookies** storing **entries**. To be more precise, according to bookkeeper's nomenclature, a **bookie** is an individual storage server while **entries** represent stored data, therefore an ensemble of size E represents simply a group of storage servers.

The aim of this design is to guarantee **consistency** in an ensemble of bookies of all stored data exploiting a **quorum-based replicated-write** protocol. As known, to support replicated writes at multiple replicas of a file, a client must first contact at least half the servers plus one (a majority) and get them to agree to do the

¹<https://github.com/apache/bookkeeper>

²<https://github.com/apache/openjpa>

³<https://www.cs.waikato.ac.nz/~ml/weka/>

⁴`org.apache.bookkeeper.client.DefaultEnsemblePlacementPolicy`

⁵`org.apache.bookkeeper.util.DiskChecker`

update[7]. Technically, in order to modify a file, a client needs to assemble the so called **write quorum**, that is an arbitrary collection of servers which must be more than the half of all available servers[7]. Therefore, using bookkeeper's nomenclature, the size of write quorum Q_w represents the number of bookies where each entry is written.

According to bookKeeper protocol[3], the following invariant must hold:

$$E \geq Q_w \geq Q_a \quad (1)$$

In other word, the ensemble size E must be larger than the write quorum size Q_w , which must in turn be larger than the so called **ack quorum size** Q_a , which represents, instead, the number of nodes an entry must be acknowledged on.

BookKeeper uses several algorithms to select a number of bookies from a cluster to build an ensemble compliant to above specifications, some of which are capable to exploit several network topologies too.

According to BookKeeper design, the implementations of these algorithm must be compliant to EnsemblePlacementPolicy interface [4]; in other words, any implementation must respect a specific *contract*, established by aforementioned interface, which covers aspects related to initialization and bookie selection for data placement and reads[4].

Currently there are 3 implementations available by default. They are:

- DefaultEnsemblePlacementPolicy
- RackawareEnsemblePlacementPolicy
- RegionAwareEnsemblePlacementPolicy

In particular, DefaultEnsemblePlacementPolicy class encapsulates the simplest algorithm for bookie selection for ensemble creation because it, simply, picks bookies randomly in order to build an ensemble.

2.2 DiskChecker class testing analysis

According to BookKeeper's documentation, DiskChecker class is used to provide several utility functions for checking disk problems, managing all directories belonging to **ledgers**, which represents the basic unit of storage in BookKeeper[1].

Specifications establishes that, for each ledger directory, is possible to specify a maximum disk space D which can be used and a, so-called, warning threshold D_w for disk usage[2]. In particular, the following invariant must hold:

$$\begin{aligned} 0 < D_w < D \\ D_w &\leq D \end{aligned} \quad (2)$$

However, for both thresholds, BookKeeper's specifications establish a default value, which is equal to 0.95.

Using out test set T_1 , which is retrievable from commit bdf2fc28d26⁶, our testing activity reports that up to 5 unit tests have failed, out of a total of 31, revealing the presence of several bugs, despite its stability according to our prediction model. In following subsections, we will describe how we have build our test sets.

⁶<https://github.com/AndreaG93/ISW2-bookkeeper/tree/bdf2fc28d216b0adf9e467cfb51c4f334e84d5c9>

2.2.1 DiskChecker(float, float). Obviously, that method represent the constructor of DiskChecker class, which manage directories belonging to ledger entities.

Using the previously given rules, we get $1 \times 1 = 1$ valid test cases (by combining the representatives of the valid equivalence classes) and $1 + 3 = 4$ negative test cases (by separately testing representatives of every invalid class). Since test cases belonging to both iEC_3 of threshold parameter and iEC_1 of warnThreshold overlap, in total 4 test cases result from the 6 equivalence classes.

2.2.2 checkDir(File dir). JaCoCo⁷

2.3 Adequacy Criteria

2.3.1 Statement coverage. Our test set T is able to cover up to 85 statements out of a total of 89, reaching, according to Sonarcloud report, a *statement coverage* equal to **0,955 (95,5%)**.

However we believe that DiskChecker class contains several statements that we can consider as **unreachable** because they fall on an *infeasible path*, that is a path that would never be reached by our test set T with any type of input data; to be more precise, we believe that setDiskSpaceThreshold(float, float) method contains up to 4 unreachable statements.

Generally, in order to identify unreachable statements, drawing the flow-graph of the developed code and finding out the path that would never be reached is required, however we can consider as unreachable all statements of setDiskSpaceThreshold(float, float) method owing to following reasons:

- (1) setDiskSpaceThreshold(float, float) method is **never used** in the project, that is it is never called by any other method.
- (2) setDiskSpaceThreshold(float, float) method has **no access modifier** (*package private*) which means, according to Java language specification, that it is only accessible within classes in the same package, therefore that method is not visible by any test set (we are assuming that test code, generally located into /test directory, is always included into a different package respect to application code, which is conversely located into /main directory)

Therefore, we can conclude by stating that:

$$\text{Statement Coverage} = \frac{|S_c|}{|S_e| - |S_i|} = \frac{85}{89 - 4} = 1 = 100\% \quad (3)$$

Since statement coverage of T is 1, we can consider T as **adequate with respect to the statement coverage criterion**.

2.3.2 Decision coverage (Branch decision coverage). According to [6], a *decision* is considered covered if the flow of control has been diverted to all possible destinations that correspond to this decision, i.e. all outcomes of the decision have been taken. This implies that, for example, the expression in the if or while statement has evaluated to true in some execution and to false in the same or another execution. Note that each if and each while contribute to *one* decision whereas a switch may contribute to more than one.

According to our analysis, DiskChecker class contains 14 possible decisions, therefore $|D_e| = 14$ where D_e is the set of decisions

⁷<https://github.com/jacoco/jacoco>

Table 1: Equivalence classes and representatives of DiskChecker method

Parameter	Equivalence Classes	Representative
threshold	$vEC_1: 0 < \text{warnThreshold} \leq x < 1$	0.5
	$iEC_1: x \leq 0$	0
	$iEC_2: x \geq 1$	1
	$iEC_3: 0 < x < \text{warnThreshold} < 1$	0.5 (while warnThreshold is 0.6)
warnThreshold	$vEC_1: x \leq \text{threshold}$	0.5 (while threshold is 0.5)
	$iEC_1: x > \text{threshold}$	0.6 (while threshold is 0.5)

Table 2: Test cases of DiskChecker method

Test Case	Parameter		Expected output	Actual output	Passed
	threshold	warnThreshold			
<i>Valid</i> ₁	0.95	0.95	No Exception	No Exception	✓
<i>Invalid</i> ₁	0	0.5	Exception	Exception	✓
<i>Invalid</i> ₂	1	0.5	Exception	Exception	✓
<i>Invalid</i> ₃	0.5	0.6	Exception	Exception	✓

Table 3: Equivalence classes and representatives of checkDir method

Parameter	Equivalence Classes	Representatives
dir	vEC_1 : Valid File object representing a valid directory (existent, readable and writeable).	<code>IOUtils.createTempDir("directoryFile", "test")</code>
	vEC_2 : Valid File object representing a not existent, but makable, directory.	<code>new File("./makeMultiple/dir/path")</code>
	iEC_1 : Valid File object which does not represent a directory, that is it can be a regular file, a symbolic link, a character device file etc.	<code>new File("/dev/zero")</code>
	iEC_2 : Valid File object representing a non-existent and non-makable directory.	<code>new File("/root/notMakable")</code>
	iEC_3 : Valid File object representing an existent, not readable directory.	<code>new File("/root")</code>
	iEC_4 : Valid File object representing an existent, not writeable directory.	<code>new File("/home")</code>
	iEC_5 : Valid File object representing an invalid directory (from file system point of view because, for example, it contains forbidden characters).	<code>new File("\u0000")</code>
	iEC_6 : A null object.	<code>null</code>

in the program. Our test set T can cover all decision, therefore $|D_c| = 14$, where D_c is the set of decisions covered.

$$\text{Decision Coverage} = \frac{|D_c|}{|D_e| - |D_i|} = \frac{14}{14 - 0} = 1 = 100\% \quad (4)$$

Since decision coverage of T is 1, we can consider T as **adequate with respect to the decision coverage criterion**.

2.3.3 Condition coverage. Unlike decision coverage, condition coverage ensures that each simple condition within a compound condition has assumed both values true and false.

Let be given:

C_e the set of simple conditions in the program.

C_c the set of simple conditions covered by our test set T .

C_i the set infeasible simple conditions.

According to JaCoCo and sonarcloud reports, our test set T covers up to 46 simple conditions, out of a total of 46, while $|C_i| = 0$. Therefore we can say that:

$$\text{Condition Coverage} = \frac{|C_c|}{|C_e| - |C_i|} = \frac{46}{46 - 0} = 1 = 100\% \quad (5)$$

Table 4: Test cases of checkDir method

Test Case	Parameter dir	Expected output	Actual output	Passed
<i>Valid₁</i>	IOUtils.createTempDir("directoryFile", "test")	No Exception	No Exception	✓
<i>Valid₂</i>	new File("./makeMultiple/dir/path")	No Exception	No Exception	✓
<i>Invalid₁</i>	new File("/dev/zero")	Exception	Exception	✓
<i>Invalid₂</i>	new File("/root/notMakable")	Exception	Exception	✓
<i>Invalid₃</i>	new File("/root")	Exception	Exception	✓
<i>Invalid₄</i>	new File("/home")	Exception	Exception	✓
<i>Invalid₅</i>	new File("\u0000")	Exception	Exception	✓
<i>Invalid₆</i>	null	Exception	Exception	✓

Table 5: Equivalence classes and representatives of getTotalFreeSpace and getTotalDiskSpace methods

Parameter	Equivalence Classes	Representatives
dir	<i>vEC₁</i> : A valid non-empty List<File> object containing File objects every of which represents a valid (existent and readable) directory.	See the code
	<i>vEC₂</i> : A valid empty List<File> object.	new ArrayList<>()
	<i>iEC₁</i> : A valid non-empty List<File> object in which at least one File object does not represent a directory, that is it can represent regular file, symbolic link, character device file etc.	See the code
	<i>iEC₂</i> : A valid non-empty List<File> object in which at least one File object represent a not existent directory.	See the code
	<i>iEC₃</i> : A valid non-empty List<File> object in which at least one File object represent an existent but not readable directory.	See the code
	<i>iEC₃</i> : A valid non-empty List<File> object in which at least one File object represent an invalid directory (from file system point of view).	See the code
	<i>iEC₃</i> : A valid non-empty List<File> object in which at least one File object is null.	See the code
	<i>iEC₅</i> : A null object.	null

Table 6: Equivalence classes and representatives of getTotalDiskUsage, getTotalFreeSpace, getTotalDiskSpace methods

Parameter	Equivalence Classes	Representatives
dir	<i>vEC₁</i> : A valid non-empty List<File> object containing File objects every of which represents a valid (existent, readable and writeable) directory.	-
	<i>iEC₁</i> : A valid non-empty List<File> object containing File objects every of which represents a not existent directory.	-
	<i>iEC₂</i> : A valid non-empty List<File> object containing File objects every of which represents an existent, not readable and not writeable directory.	-
	<i>iEC₃</i> : A valid non-empty List<File> object containing File objects every of which not represents a directory, that is it can represent regular file, symbolic link, character device file etc.	-
	<i>iEC₄</i> : A valid empty List<File> object.	new ArrayList<>()
	<i>iEC₅</i> : A null object.	null

Since condition coverage is equal to 1, our test set T is adequate with respect to the condition coverage criterion.

2.4 Mutation Analysis

According to PIT ⁸ report, mutation coverage equal to 63, con 46

2.5 Equivalence Class Partitioning

⁸<https://pitest.org/>

Listing 1: onClusterChanged method signature.

2.5.1 onClusterChanged.

```
Set<BookieSocketAddress> onClusterChanged(Set<BookieSocketAddress> writableBookies, Set<BookieSocketAddress> readOnlyBookies, int ensembleSize, int writeQuorumSize, int ackQuorumSize, Map<String, String> customMetadata, boolean excludeBookies)
```

Table 7: Equivalence classes and representatives of onClusterChanged method

Parameter	Equivalence Classes	Representatives
writableBookies	EC_1 : Valid Set<BookieSocketAddress> object (size ≥ 0) containing valid BookieSocketAddress objects. EC_2 : null object	<p>The implementation should respect to the replace settings. The size of the returned bookies list should be equal to the provided ensembleSize.</p> <p>The @NotNull Annotation is, actually, an explicit contract declaring the following: A method should not return null. A variable (like fields, local variables, and parameters) cannot hold null value.</p>
readOnlyBookies	EC_1 : Valid Set<BookieSocketAddress> object (size ≥ 0) containing null items or invalid BookieSocketAddress objects. EC_2 : null object	<p>Morevoer, Choose numBookies bookies for ensemble. If the count is more than the number of available nodes, BKException.BKNotEnoughBookiesException is thrown.</p> <p>The implementation should respect to the replace settings. The size of the returned bookies list should be equal to the provided ensembleSize.</p> <p>The @NotNull Annotation is, actually, an explicit contract declaring the following: A method should not return null. A variable (like fields, local variables, and parameters) cannot hold null value.</p>

Table 8: Test cases of onClusterChanged method

Test Case	Parameter	Expected output	Actual output	Passed
Valid, writableBookies	Valid Set<BookieSocketAddress> object (size = 0).	Empty Set<BookieSocketAddress>	Empty Set<BookieSocketAddress>	✓
Valid, writableBookies	Valid Set<BookieSocketAddress> object inside which there is 1 valid BookieSocketAddress object.	Empty Set<BookieSocketAddress>	Empty Set<BookieSocketAddress>	✓
Valid, writableBookies	Valid Set<BookieSocketAddress> object inside which there is 1 valid BookieSocketAddress object.	Not empty Set<BookieSocketAddress> (size = 4)	Not empty Set<BookieSocketAddress> (size = 4)	✓
Invalid, writableBookies	Valid Set<BookieSocketAddress> object inside which there is one null reference.	Exception	Empty Set<BookieSocketAddress>	✗
Invalid, writableBookies	Valid Set<BookieSocketAddress> object inside which there is 1 valid BookieSocketAddress object.	Exception	Empty Set<BookieSocketAddress>	✗
Invalid, writableBookies	null	Exception	Exception	✓
Invalid, writableBookies	Valid Set<BookieSocketAddress> object inside which there is 1 valid BookieSocketAddress object.	Exception	Exception	✓

2.5.2 newEnsemble.

```
public void updateBookieInfo(Map<BookieSocketAddress, BookieInfo> bookieInfoMap)
```

According to bookkeeper documentation [4], newEnsemble method takes up to 5 parameters and is used to build an ensemble made up of several bookies.

ensembleSize Represents the ensemble size.

writeQuorumSize the value of write quorum size

ackQuorumSize the value of ack Quorum Size

customMetadata user added metadata. This data are not managed directly by DefaultEnsemblePlacementPolicy class

excludeBookies ookies that should not be considered as targets for the new ensemble.

is used to build an ensemble made up ensembleSize bookies

which takes up to 5 parameters:

According to method signature, if not enough bookies are available to build an ensemble of size ensembleSize, a BKNotEnoughBookiesException is thrown.

. Moreover, Choose numBookies bookies for ensemble. If the count is more than the number of available nodes, BKException.BKNotEnoughBookiesException is thrown.

. Moreover,

Choose numBookies bookies for ensemble. If the count is more than the number of available nodes, BKException.BKNotEnoughBookiesException is thrown.

The implementation should respect to the replace settings. The size of the returned bookies list should be equal to the provided ensembleSize.

The @NotNull Annotation is, actually, an explicit contract declaring the following: A method should not return null. A variable (like fields, local variables, and parameters) cannot hold null value.

A method should not return null. A variable (like fields, local variables, and parameters) cannot hold null value.

2.5.3 replaceBookie.

2.5.4 updateBookieInfo. According to DefaultEnsemblePlacementPolicy

interface's documentation [7], updateBookieInfo is used to update bookie info details taking only one input parameter, a Map<BookieSocketAddress, BookieInfo> object. Method signature is the following:

```
public void updateBookieInfo(Map<BookieSocketAddress, BookieInfo> bookieInfoMap)
```

Listing 2: onClusterChanged method signature.

```
public void updateBookieInfo(Map<BookieSocketAddress, BookieInfo> bookieInfoMap)
```

REFERENCES

- [1] [n.d.]. BookKeeper concepts and architecture. <https://bookkeeper.apache.org/docs/4.8.2/getting-started/concepts/>. [Online; accessed 4-October-2020].
- [2] [n.d.]. BookKeeper configuration. <https://bookkeeper.apache.org/docs/4.7.3/reference/config/>. [Online; accessed 6-October-2020].
- [3] [n.d.]. The BookKeeper protocol. <https://bookkeeper.apache.org/docs/4.8.2/development/protocol/>. [Online; accessed 5-October-2020].
- [4] [n.d.]. Interface EnsemblePlacementPolicy. <https://bookkeeper.apache.org/docs/4.8.2/api/javadoc/org/apache/bookkeeper/client/EnsemblePlacementPolicy.html>. [Online; accessed 5-October-2020].
- [5] Aalok Ahluwalia, Massimiliano Di Penta, and Davide Falessi. 2020. On the Need of Removing Last Releases of Data When Using or Validating Defect Prediction Models.
- [6] Andreas Spillner, Tilo Linz, and Hans Schaefer. 2011. *Software Testing Foundations: A Study Guide for the Certified Tester Exam* (3rd ed.). Rocky Nook.
- [7] M. van Steen and A.S. Tanenbaum. 2017. *Distributed Systems, 3rd ed.* distributed-systems.net.

Table 9: Equivalence classes and representatives of newEnsemble method

Parameter	Equivalence Classes	Representatives
ensembleSize	vEC_1 : A not null Map<BookieSocketAddress, BookieInfo> object (size ≥ 0) containing valid (BookieSocketAddress, BookieInfo) objects pairs. In particular, for any BookieSocketAddress object is true that:	
	vEC_1 : A not null Map<BookieSocketAddress, BookieInfo> object (size ≥ 0) containing valid (BookieSocketAddress, BookieInfo) objects pairs. In particular, for any BookieSocketAddress object is true that:	
	vEC_1 : A not null Map<BookieSocketAddress, BookieInfo> object (size ≥ 0) containing valid (BookieSocketAddress, BookieInfo) objects pairs. In particular, for any BookieSocketAddress object is true that:	
	vEC_1 : A not null Map<BookieSocketAddress, BookieInfo> object (size ≥ 0) containing valid (BookieSocketAddress, BookieInfo) objects pairs. In particular, for any BookieSocketAddress object is true that:	
writeQuorumSize	vEC_1 : A not null Map<BookieSocketAddress, BookieInfo> object (size ≥ 0) containing valid (BookieSocketAddress, BookieInfo) objects pairs. In particular, for any BookieSocketAddress object is true that:	
ackQuorumSize	vEC_1 : A not null Map<BookieSocketAddress, BookieInfo> object (size ≥ 0) containing valid (BookieSocketAddress, BookieInfo) objects pairs. In particular, for any BookieSocketAddress object is true that:	

Table 10: Equivalence classes and representatives of updateBookieInfo method

Parameter	Equivalence Classes	Representatives
customMetadata		
excludeBookies		

Table 11: Equivalence classes and representatives of updateBookieInfo method

Parameter	Equivalence Classes	Representatives
bookieInfoMap	<p>A not null Map<BookieSocketAddress, BookieInfo> object (size ≥ 0) containing valid (BookieSocketAddress, BookieInfo) objects pairs. In particular, for any BookieSocketAddress object is true that:</p> <p>vEC_1: (1) port field is ≥ 0; (2) hostname field is not null;</p> <p>Moreover, for all BookieInfo object is true that:</p> <p>(1) totalDiskSpace field is ≥ 0; (2) freeDiskSpace field is ≥ 0; (3) freeDiskSpace \leq totalDiskSpace;</p>	(see the code)
	<p>iEC_1: A not null Map<BookieSocketAddress, BookieInfo> object (size ≥ 0) in which at least one invalid (BookieSocketAddress, BookieInfo) pair object is present; for example, there is one pair in the form (BookieSocketAddress, null), (null, BookieInfo) or (null, null).</p>	(see the code)
	<p>iEC_2: A not null Map<BookieSocketAddress, BookieInfo> object (size ≥ 0) in which at least one malformed (BookieSocketAddress, BookieInfo) pair object is present. In particular, there is at least one BookieSocketAddress in which:</p> <p>(1) port field is < 0. (2) hostname field is null.</p> <p>Moreover, there is at least one BookieInfo in which:</p> <p>(1) totalDiskSpace field is < 0 (2) freeDiskSpace field is < 0 (3) freeDiskSpace $>$ totalDiskSpace;</p>	(see the code)
	iEC_3 : null object.	null

Table 12: Test cases of updateBookieInfo method

Test Case	Parameter	Expected output	Actual output
	bookieInfoMap		
$Valid_1$	A not null Map<BookieSocketAddress, BookieInfo> object (size = 0).	Any Exception	Any Exception
$Valid_2$	A not null Map<BookieSocketAddress, BookieInfo> object (size ≥ 0) containing valid (BookieSocketAddress, BookieInfo) objects pairs.	Any Exception	Any Exception
$Valid_2$	A not null Map<BookieSocketAddress, BookieInfo> object (size ≥ 0) containing valid (BookieSocketAddress, BookieInfo) objects pairs. The isWeighted field of DefaultEnsemblePlacementPolicy object is set to false	Any Exception	Exception